

How is the web written (and how to read it)

SICSS

 CREST



INSTITUT
POLYTECHNIQUE
DE PARIS

Outline

1. The WWW is not the Internet

- A Brief History of the Internet
- Birth of the WWW
- Internet, HTML and Browsers

2. The WWW, Back & Front

- How is the web written?
- The structure of a simple webpage
- The *www millefeuille*

3. How-to in R

4. A note on APIs

5. Headless browsers

The WWW is not the Internet

1. A Brief History of the Internet

The WWW is not the Internet

1. A Brief History of the Internet

- *Competing Histories*
- A “Cold War Technology”
 - DARPA and the race for the future of technology
 - DARPA & ARPA-NET (1958-1969)




Fig. 1—(a) Centralized. (b) Decentralized. (c) Distributed networks.

Безопасное путешествие лайка


Paul Baran, 1962

The WWW is not the Internet

1. A Brief History of the Internet

- *Competing Histories*

- “The Hippies did it”: Libertarian origins
 - Augmenting individuals: D. Engelbart against technoscience
 - Collective collaboration and the rise of hackers: science (and connected computers) for the people.



The WWW is not the Internet

1. A Brief History of the Internet

- *Technological & Political Struggles*
 - From circuit switching to packet switching
 - A flurry of networks → TCP/IP (1978-1983)
 - Whose Technology?
 - 1980s: NSF grant to connect US universities

The French Contender (the Minitel)





Fig. 1—(a) Centralized. (b) Decentralized. (c) Distributed networks.



The WWW is not the Internet

2. Birth of the WWW

- *A network that uses the internet*
 - Tim Berners-Lee (1989)
- *A decentralized sorting system*
 - Child of the previous evolution
 - The hyperlink at its core
- *HTML & HTTP*
 - HTML: Hyper Text Markup Language
 - HTTP: Hyper Text Transfer Protocol



The WWW is not the Internet

2. Birth of the WWW

- *The Internet as a Product*
- IMDB 1990; Amazon, Ebay, Craigslist 1995;
- Hotmail 1996; Yahoo, Google, Paypal 1998; 2001 Wiki;etc

Results Summary			
Month	# of Web sites	% .com sites	Hosts* per Web server
6/93	130	1.5	13,000 (3,846)
12/93	623	4.6	3,475 (963)
6/94	2,738	13.5	1,095 (255)
12/94	10,022	18.3	451 (99)
6/95	23,500	31.3	270 (46)
1/96	100,000	50.0	94 (17)
6/96	230,000 (est)	68.0	41
1/97	650,000 (est)	62.6	NA



Avenue Q: The internet is for...

The WWW is not the Internet

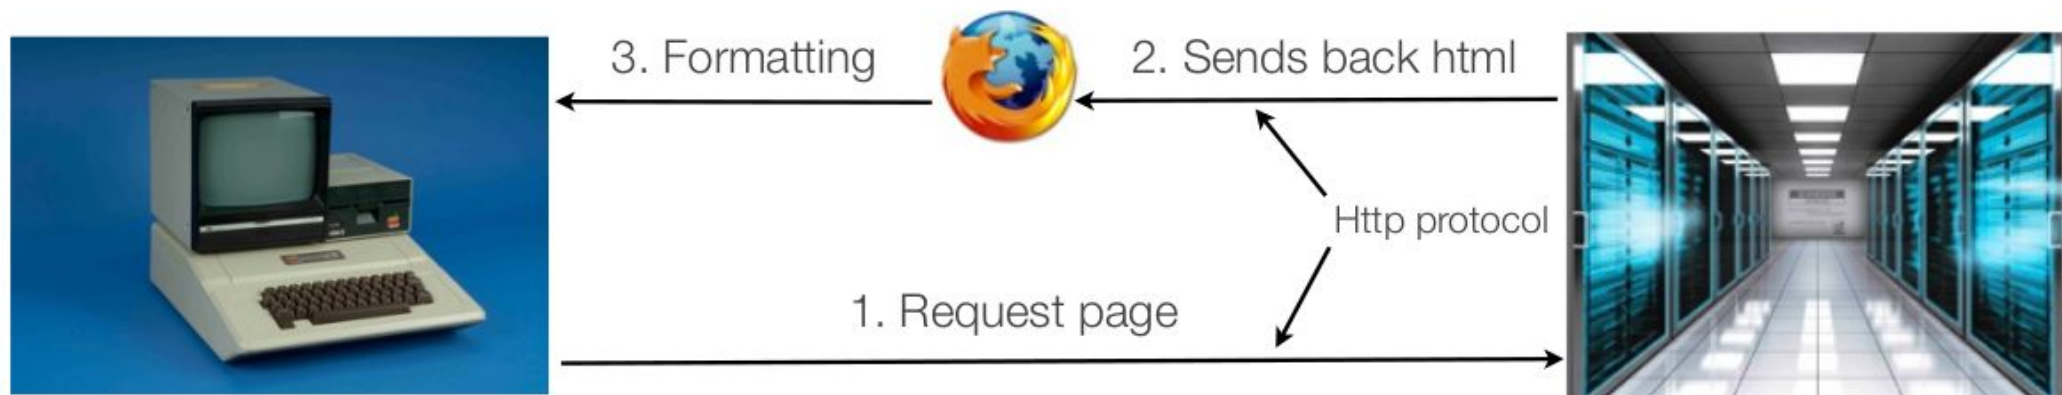
3. The Internet, the WWW and Browsers

- . Browsers were essential to popularizing the WWW*

The WWW is not the Internet

3. The Internet, the WWW and Browsers

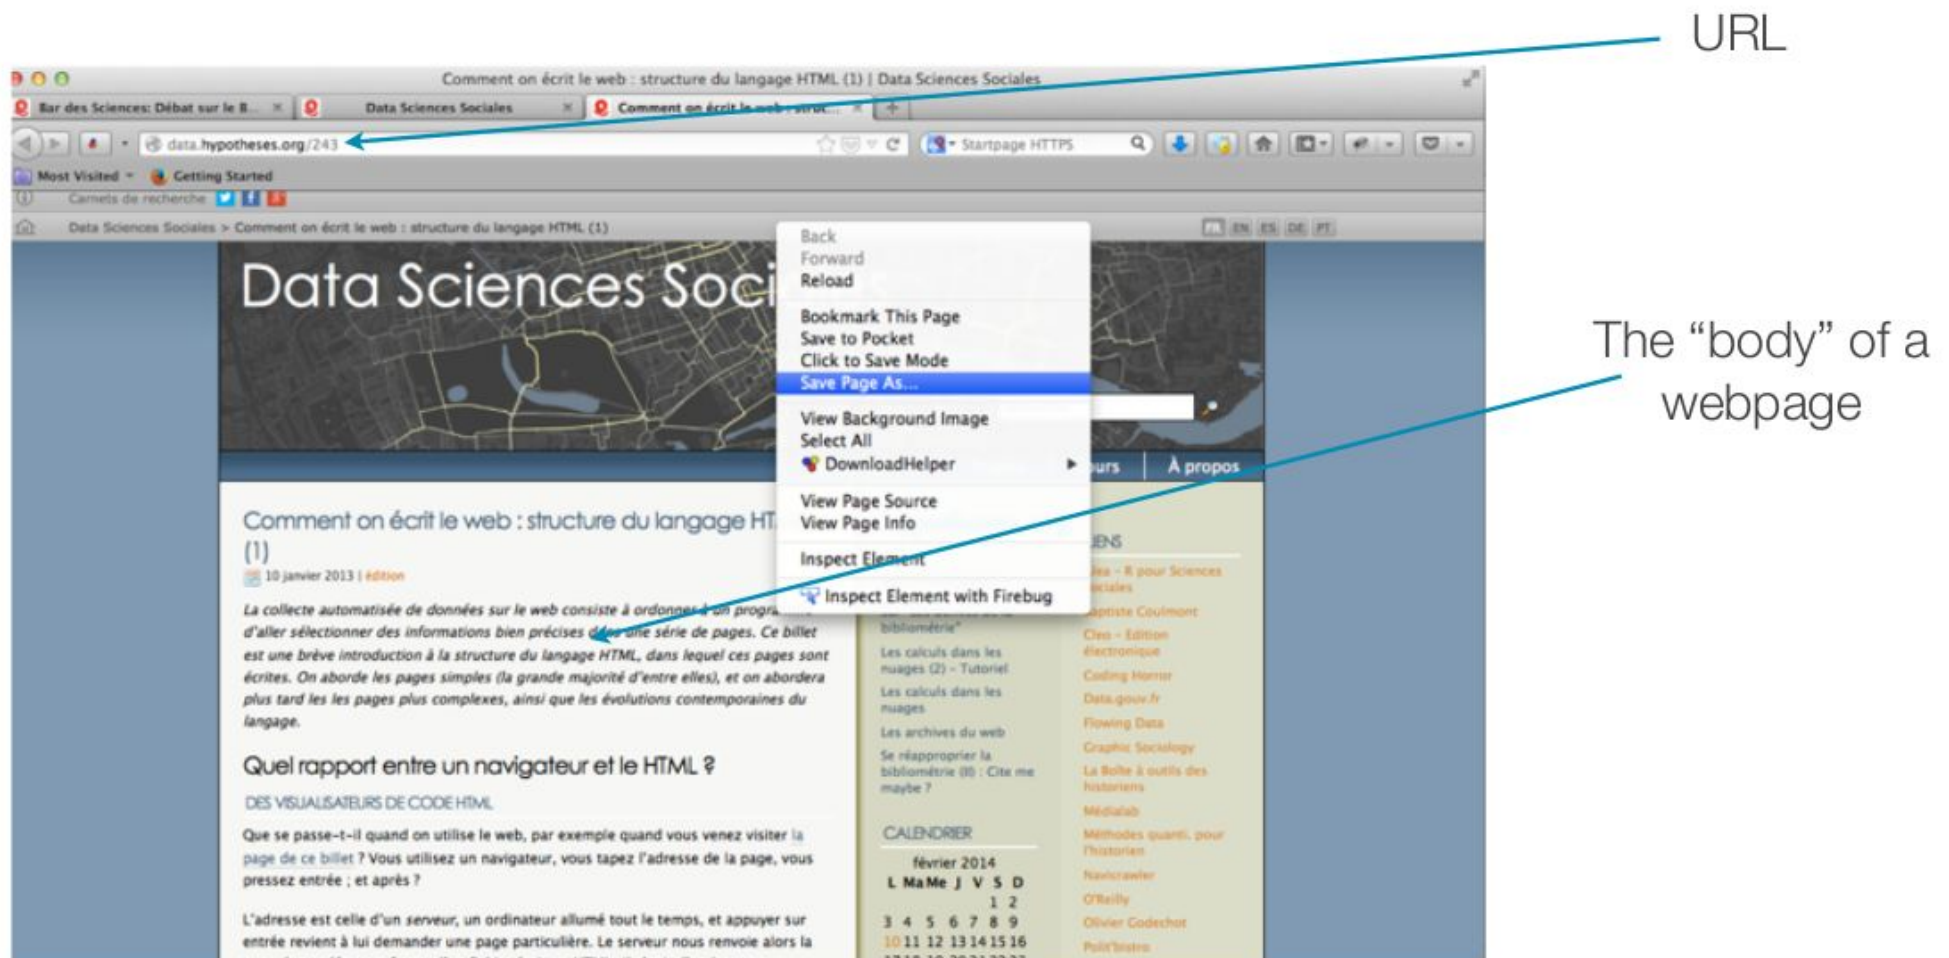
- *Browsers send, receive and interpret code*
- The web relies on the circulation of text in HTML
- The **www** is based on communication between computers via a protocol called **HTTP**
- Computers & pages are identified by their address, called a **URL** (Uniform Resource Locator)
- **HTML** files are transferred and subsequently formatted into a legible format.



The WWW is not the Internet

3. The Internet, the WWW and Browsers

- *Browsers visualize code*



The WWW, Back & Front

The WWW, Back & Front

1. How is the Web Written?


- *The Web is Premised on HTML*
 - A “Rich Language”
 - A Structuring Principle: tags
 - HTML works with tags
 - The text displayed is surrounded by extra information, contained in these containers.
- Ex. <TAG> This is very interesting </TAG>

The WWW, Back & Front

1. How is the Web Written?

- *The Web is Premised on HTML*
 - A “Rich Language” (more than meets the eye)
 - A Structuring Principle: **tags**
 - Tags have a **type** and an **attribute**
 - *Types* are fix (a, span, li, div). They have a limited number of attributes.
 - For more, see [this page](#).

 This is very interesting



The WWW, Back & Front

1. How is the Web Written?

. *Head & Body*

```
<!DOCTYPE html ...>
<html>
  <head>
    <meta ... >
    <link ...>
    <script ...>
    ...
  </head>
  <body>
    ...
  </body>
</html>
```

← Type of document (declares format)

← Head (meta informations about the document)

← Body of text (where most of the action happens for us)

The WWW, Back & Front

1. How is the Web Written?

- *A few common tags*
 - `<div>` : block of text
 - `<p>` : paragraph
 - `<a>` : hypertext link
 - `<h1>` : (resp. h2, h3, h4, h5) titles
 - `<!...>` : Comment

The WWW, Back & Front

2. The Structure of a Simple Webpage

. *What happens in the code is visible on the page*

Display source code

The screenshot shows a web browser window with the URL `data.hypotheses.org/243`. The page title is 'Comment on écrit le web : structure du langage HTML (1) | Data Sciences Sociales'. The page content includes a header with the title 'Data Sciences Sociales', a main text area, and a sidebar with a calendar and a list of links. A context menu is open over the page, with the 'View Page Source' option selected. An arrow points from the text 'Display source code' to the 'View Page Source' option in the menu.

The source code is displayed in a separate window titled 'Source of: http://data.hypotheses.org/page/2'. The code is HTML and includes comments in French. The comments explain the structure of HTML, including the DOCTYPE declaration, the head section, and the body section. The code also includes a link to the CSS file and a meta tag for the page title.

The WWW, Back & Front

2. The Structure of a Simple Webpage

- Tools in your browsers help you “inspect elements”*

The screenshot displays a web browser window with a webpage for Étienne Ollion. The page features a header with the name "Étienne Ollion" and his title "Chercheur CNRS | Professeur associé à l'Ecole polytechnique". Below the header is a navigation menu with links: ACCUEIL, COURS, PUBLICATIONS, RECENT, and ENGLISH. The main content area is titled "Accueil" and contains a large image of a theater interior with red seats and a stage. To the right of the image is a calendar for April 2019.

The browser's developer tools are open at the bottom, showing the HTML structure of the page. The selected element is the `<body>` tag, which has the class `home page-template-default page page-id-7 wp-embed-responsive custom-font-enabled single-author`. The HTML structure is as follows:

```
<!DOCTYPE html>
<!--[if IE 7]> <html class="ie ie7" lang="fr-FR"> <![endif-->
<!--[if IE 8]> <html class="ie ie8" lang="fr-FR"> <![endif-->
<!--[if !(IE 7) & !(IE 8)]><!-->
<html lang="fr-FR">
  <!--[endif]-->
  <head>
    <!--[endif]-->
  </head>
  <body class="home page-template-default page page-id-7 wp-embed-responsive custom-font-enabled single-author">
    <div id="page" class="hfeed site">
      <header id="masthead" class="site-header" role="banner">
        <!--[endif]-->
      </header>
    </div>
  </body>
</html>
```

The right sidebar of the developer tools shows the "Layout" tab, which displays the box model for the selected element. The box model is defined by the following CSS rules:

- `flex-grow: 1;`
- `flex-shrink: 1;`
- `flex-basis: 0px;`
- `margin: 0px;`

The WWW, Back & Front

3. The *www millefeuille*

- A webpage is built on HTML
- And it includes other types of files
 - Pictures
 - Content Style Sheet (CSS)
 - Javascript

> Increasingly, the web has become a *millefeuille*

The WWW, Back & Front

3. The *www millefeuille*

> Increasingly, the web has become a *millefeuille*

This has consequences for scraping. But keep in mind that a regularity on the screen means regularity in the code. We are going to use this

How-to in R

1. Finding your way around

There is a wealth of dedicated libraries

This page maintains a list of all that there is at moment
(and it's plenty)

Scraping: *httr* or ***rvest***

Selecting in HTML (or XML): *XML* or ***rvest***

Selection in json: *rjson*, *rjsonio*, *jsonlite*...

How-to in R

2. Basic instructions

`read_html()` will read the page and transform it into an XML document.

Thus,

`read_html("https://sicss.io/2022/paris/schedule")` will output the source code of the schedule page for the SICSS-Paris program.

How-to in R

2. Basic instructions



How-to in R

2. Basic instructions

Yes, in ~70% of the cases, all you need to do to scrape a page is to do

```
read_html("PAGE")
```

How-to in R

Sometimes that won't be enough, **and the website will see the evil crawler in you.**

You will need to dress-up like an honest browser

This goes through the `user_agent` command (*httr*)

```
user_agent("Mozilla/5.0 (Macintosh;U;  
Intel MacOS X 10.6; en-US")
```

```
user_agent("roger.rabbit@gmail.com")
```



How-to in R

Sometimes that won't be enough, **because you'll need cookies.**

With rvest, you will have to create a **session**, which stores the said cookies and allows you to navigate from there.

And then you'll use **read_html()**

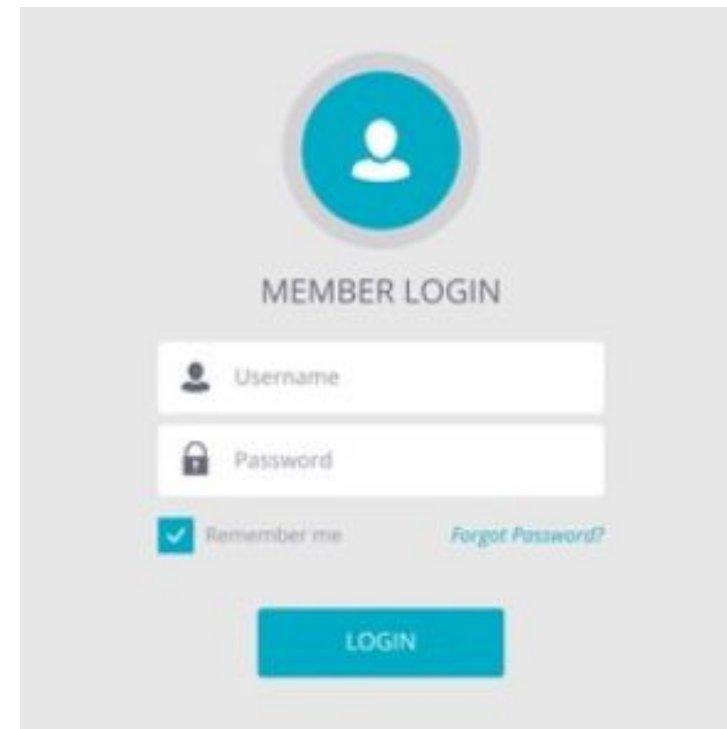


How-to in R

Sometimes that won't be enough, **because you'll need to log in.**

You'll use a function called **html_form()**,

And then you'll use **read_html()**



MEMBER LOGIN

Username

Password

☒ Remember me [Forgot Password?](#)

LOGIN

How-to in R

Once you have done that:

Great news: we are back to square 1!

(You could have copied and pasted the source code in your console, couldn't you?)

Except that now, it is in R, and you can use it

A quick note on APIs

APIs


Application Programmers Interface

- A feature of the web 2.0
- Legal, and often easier
- Different forms: login, with rate limits, or just open

APIs

A simple yet elaborate API

- <https://www.openstreetmap.org/search?query=ecole%20polytechnique%20paris>



APIs

Most APIs require registration

See the now ~~over~~often-used Twitter API for academics

1. Register for a project
<https://apps.twitter.com>
2. Wait for approval
3. Get your credentials
4. Make requests

See Chris Bail's [detailed page](#) on APIs

See [academicTwitterR](#) webpage

The screenshot displays the Twitter Developer Portal interface. On the left is a dark sidebar with the 'Developer Portal' logo and navigation links: 'Dashboard', 'Projects & Apps', 'Products' (marked as 'NEW'), and 'Account'. The main content area shows the 'ACADEMIC PROJECT' setup for 'Politicians 2.0'. It includes tabs for 'Overview' and 'Settings'. Under the 'Access' section, there is a table for 'Academic research' with the following details:

Academic research	
Apps	1 environment per project
Tweets	10M Tweets per month / Project
Cost	free
License	For non-commercial use only

Below this, the 'Usage' section shows a 'MONTHLY TWEET CAP USAGE' progress bar at 0%, indicating '0 Tweets pulled of 10,000,000'. It also notes that the usage resets on April 13 at 00:00 UTC. At the bottom of the page, there are links for 'PRIVACY', 'COOKIES', 'TWITTER TERMS & CONDITIONS', 'DEVELOPER POLICY & TERMS', and a copyright notice for © 2022 TWITTER INC.

APIs

Back to legal and deontological matters

- If there is an API, use it
- If there is no API, ask yourself: are you doing something illegal?
- Sure, scraping Twitter is legal, and its content is public. But what do you learn from individuals? And how should you protect them?


A Quick Note on Headless Browsers

A growing trend in the web industry is to have websites that respond to your behavior (scrolling, clicking, etc).



A Quick Note on Headless Browsers


For behavior-based websites



Sön 18 apr kl 14:30

Socialt planerad sekelskiftesdröm renoverad till perfektion med två fungerande eldstäder på Kungsholmens bästa läge. Här möts ni av en

Plus



Västmannagatan 70, VÅN 3

Vasastan, Stockholm

7 995 000 kr
2 787 kr/mån

72 m²
111 042 kr/m²

3 rum

Välkommen till denna genomgående sekelskiftes-3a högt upp i huset

Balkong Hiss

UNIK
FASTIGHETSFORMIDLING

Visar 1 - 50 av 2428

1

2

3

...

49

Nästa »

A Quick Note on Headless Browsers

To do this, you will need to use Javascript in order to create a “**headless browser**”, i.e. a browser piloted from your command line.

A Quick Note on Headless Browsers

To do this, you will need to use Javascript in order to create a “**headless browser**”, i.e. a browser piloted from your command line.

This is slightly more complex as you need to install other software, but we'll see an example later.

This is also an easy way to avoid some classic headaches.

A Quick Note on Headless Browsers

In R, this is often done using “Selenium”.
To do so, install “RSelenium”

For me, it worked better installing Docker too (...)

And you will need to type, in the command line, a few lines of code. See this explanation by Chris Bail.

Conclusion



But all you need to know is **read_html()**
(and where to look for)